ARYAN
7116951

🔷 Enterprise AI Quiz (Azure AI Foundry + Integration)

Q1.

A retail company wants to deploy a chatbot using Azure AI Foundry that integrates with their existing CRM system. What is the **best approach** to ensure seamless integration?

A. Build a standalone chatbot with no backend integration
B. Use REST APIs to connect AI services with CRM
C. Export CRM data manually to AI model
D. Use only static responses

Q2.

Your organization wants to deploy multiple AI models across departments with centralized control. Which architecture is most suitable?

A. Decentralized model deployment
B. Centralized AI platform with shared services
C. Independent model hosting per team
D. Local machine deployment

Q3.

While integrating Azure AI Foundry into an enterprise system, latency issues are observed. What is the most effective solution?

A. Reduce data size manually
B. Use caching and edge deployment
C. Stop using APIs
D. Increase UI response time

Q4.

A banking system requires strict governance and audit trails for AI models. Which feature should be prioritized?

A. Model speed
B. Model explainability and logging
C. UI design
D. Data duplication

Q5.

A company wants to retrain AI models automatically based on new data. Which approach should be implemented?

A. Manual retraining monthly
B. CI/CD pipeline with automated retraining
C. Ignore new data
D. Delete old models

Q6.

During AI deployment, data privacy concerns arise. What is the best practice?

A. Share all data openly
B. Implement role-based access control (RBAC)
C. Store data in plain text
D. Avoid encryption

Q7.

A logistics company needs real-time predictions from AI models. Which deployment pattern should be used?

A. Batch processing
B. Real-time inference endpoints
C. Manual processing
D. Offline analysis

Q8.

Your enterprise AI system must scale during peak traffic. What Azure feature helps achieve this?

A. Fixed virtual machines
B. Auto-scaling with load balancing
C. Manual scaling
D. Local deployment

Q9.

An AI solution is failing due to inconsistent data formats from multiple systems. What should be implemented?

A. Ignore inconsistencies
B. Data standardization and ETL pipelines
C. Manual correction
D. Delete data

Q10.

A company wants to monitor AI model performance over time. Which approach is best?

A. Ignore monitoring
B. Implement model monitoring and drift detection
C. Rebuild model daily
D. Disable logs

Q11.

You are integrating AI into a legacy ERP system. What is the safest approach?

A. Replace entire ERP
B. Use API-based integration layer
C. Modify core ERP code directly
D. Avoid integration

Q12.

For enterprise AI governance, which practice is most critical?

A. No documentation
B. Version control and audit logs
C. Random deployments
D. Manual tracking

Q13.

A healthcare AI solution must comply with regulations. What should be ensured?

A. Fast responses only
B. Compliance with data regulations and encryption
C. Minimal testing
D. Open access

Q14.

A company wants to reuse AI components across projects. What is the best strategy?

A. Build from scratch every time
B. Create modular and reusable AI services
C. Avoid documentation
D. Use hardcoded logic

Q15.

While deploying AI models globally, users face performance issues. What should be done?

A. Use single region deployment
B. Deploy across multiple regions (geo-distribution)
C. Stop deployment
D. Reduce users

Q16.

An enterprise wants to integrate AI with real-time dashboards. Which integration method is best?

A. Manual data updates
B. Event-driven architecture with streaming
C. Static reports
D. Batch uploads

Q17.

During scaling, cost optimization becomes critical. What should be used?

A. Always-on resources
B. Pay-as-you-go and auto-scaling
C. Fixed pricing only
D. Unlimited resources

Q18.

A fraud detection system must respond instantly. Which architecture is ideal?

A. Batch AI processing
B. Stream processing with real-time inference
C. Offline training only
D. Manual verification

Q19.

AI models are producing biased outputs. What is the best corrective action?

A. Ignore bias
B. Implement fairness checks and retraining
C. Remove model
D. Reduce data

Q20.

A company wants end-to-end AI lifecycle management. What should they implement?

A. Isolated tools
B. Integrated MLOps pipeline
C. Manual workflows
D. No monitoring